

SML_assignment_3

November 22, 2016

1 Exercise 1 - Bayesian Linear Regression

1.1 1.1

We begin by computing $p(t, x, \mathbf{x}, \mathbf{t})$, using the results from exercise 2 week 8:

$$p(t \mid x, \mathbf{x}, \mathbf{t}) = \mathcal{N}(t \mid m(x), s^2(x))$$

$$m(x) = \phi(x)^T \mathbf{m}_n = N\beta(1, x) \mathbf{S}_N \begin{pmatrix} \bar{\mu}_t \\ \bar{\mu}_{xt} \end{pmatrix}$$

$$s^2(x) = \beta^{-1} + \phi(x)^T \mathbf{S}_N \phi(x) = \beta^{-1} + (1, x) \mathbf{S}_N \begin{pmatrix} 1 \\ x \end{pmatrix}$$

$$\mathbf{S}_N^{-1} = \begin{pmatrix} \alpha & 0 \\ 0 & \alpha \end{pmatrix} + N\beta \begin{pmatrix} 1 & \bar{\mu}_x \\ \bar{\mu}_x & \bar{\mu}_{xx} \end{pmatrix}$$

We are given

$$x = \begin{pmatrix} .4 \\ .6 \end{pmatrix}$$

$$t = \begin{pmatrix} .05 \\ -.35 \end{pmatrix}$$

Hence, we compute :

$$\bar{\mu}_t = \frac{1}{N} \sum_n t_n = -.15$$

$$\bar{\mu}_x = \frac{1}{N} \sum_n x_n = .5$$

$$\bar{\mu}_{xt} = \frac{1}{N} \sum_n x_n t_n = -.095$$

$$\bar{\mu}_{xx} = \frac{1}{N} \sum_n x_n^2 = .26$$

Confirmation with python 3.5:

```

In [1]: # START PRE-AMBLE
import numpy as np
from pylab import *
# use latex interpreter
rc('text', usetex=True)
import scipy.stats as stats
import seaborn as sb
# END - PREAMBLE

ar = np.array; inv = np.linalg.inv
# given parameters
a = 2 ; b = 10
x = ar([0.4, .6])
t = ar([.05, -.35])

# set number of observed points
N = len(x)
# define mu parameters
mu_bar_t = np.mean(t)
mu_bar_xt = np.mean(t * x)
mu_bar_x = np.mean(x)
mu_bar_xx = np.mean(x**2)

print('mu_bar_t={0}\nmu_bar_x = {1}\nmu_bar_xt={2}\nmu_bar_xx= {3}'\
      .format(mu_bar_t, mu_bar_x, mu_bar_xt, mu_bar_xx))

mu_bar_t = -0.15
mu_bar_x = 0.5
mu_bar_xt = -0.095
mu_bar_xx = 0.26

```

Now we can compute S_n , $m(x)$, and $s^2(x)$:

$$\mathbf{S}_N^{-1} = \begin{pmatrix} 22 & 10 \\ 10 & 7.2 \end{pmatrix}$$

$$\mathbf{S}_n = \begin{pmatrix} .1233 & -.1712 \\ -.1712 & .3767 \end{pmatrix}$$

$$\begin{aligned} s^2(x) &= \beta^{-1} + (1, x) \mathbf{S}_N \begin{pmatrix} 1 \\ x \end{pmatrix} \\ &= \frac{1}{10} + (1, x) \begin{pmatrix} .1233 & -.1712 \\ -.1712 & .3767 \end{pmatrix} \begin{pmatrix} 1 \\ x \end{pmatrix} \\ &= .1 + .1233 - .1712x - .1712x + .3767x^2 \\ &= .2233 - .3425x + .3767x^2 \end{aligned}$$

$$\begin{aligned} m(x) &= N\beta(1, x) \mathbf{S}_N \begin{pmatrix} \bar{\mu}_t \\ \bar{\mu}_{xt} \end{pmatrix} \\ &= 20(1, x) \begin{pmatrix} .1233 & -.1712 \\ -.1712 & .3767 \end{pmatrix} \begin{pmatrix} -.15 \\ -.095 \end{pmatrix} \\ &= (1, x) \begin{pmatrix} -0.04452055 \\ -0.20205479 \end{pmatrix} \\ &= -0.04452055 - 0.20205479x \end{aligned}$$

1.2 1.2

Note we must use the square root of $s^2(x)$ as this would be the standard deviation. Using the results from question 1 we plot:

```
In [2]: # see previous exercise for these results:
m_x = lambda x: -0.04452055 - 0.20205479 * x
s2_x = lambda x: .2233 - .3425 * x + .3767 * x**2

# domain over [0,1]
x_domain = np.linspace(0,1)

mu = m_x(x_domain)
std = np.sqrt(s2_x(x_domain))

# %matplotlib inline
# plot results
fig = figure();
ax = fig.add_subplot(111)
ax.plot(\
    x_domain, mu, \
    'b-', label = 'mean')
# draw +- std as an area like in bisschop
ax.fill_between(x_domain, mu + std, mu - std, \
    facecolor = 'red', alpha =.2)
```

```

ax.plot(x, t, 'g.', label = 'data');
ax.set_xlabel('x');
ax.legend()
ax.set_title('Ex. 1.2');
savefig('Figures/12.png')

```

1.3 1.3

We know that $p(w | \alpha) = \mathcal{N}(w, | 0, \alpha^{-1}, I)$; From week 8 exercise 2, we see that the posterior for w is given as $p(w | t, x) = \mathcal{N}(w | m_n, S_n)$, with $m_N = \beta S_n \phi^T t$ and $S_n^{-1} = \alpha I + \beta \phi^T \phi$.

Thus I will write code that will: - Create multivariate normal object with $\mu = m_N$, $\Sigma = S_N$ - Draw 5 samples from this distribution and plot according to: $y = w_0 + w_1 * x$

```

In [3]: # load multivariate normal object
from scipy.stats import multivariate_normal as mv

phi = np.vstack((np.ones(x.shape), x))
# print(phi); assert 0
s_n = np.linalg.inv(a * np.eye(N) + b * phi.T.dot(phi))
m_n = b * s_n .dot(phi.T.dot(t))

# create object for p(w | 0, a^-1 I)

# p_w = mv([0,0], 1/a * np.eye(N))
p_w = mv(m_n, s_n)

# draw 5 random samples
random_w = p_w.rvs(5).T
# create inline plot function for the line
plot_w = lambda x, w0, w1: w0 + x * w1

# open figure and start plotting
fig = figure()
ax = fig.add_subplot(111)

# plot mean and standard deviation
# mean
ax.plot(\
    x_domain, m_x(x_domain), \
    'b-', label = 'optimal');
# standard deviation; plotting as an area
ax.fill_between(x_domain, mu + std, mu-std, \
    facecolor = 'r', alpha = .2)
# plot the data
ax.plot(x, t, 'g.', label = 'data')

# create 5 straight lines
[ax.plot(\

```

```

        x_domain, plot_w(x_domain, w[0], w[1]),\
        '--', label = 'sample {}'.format(idx)\
    )\
    for idx, w in enumerate(random_w.T)];

# plot formatting
ax.legend(loc = 0,\
        ncol = 1,\
        frameon = True,\
        framealpha = 1,\
        bbox_to_anchor = (1., 1.));

ax.grid('off');
sb.set_style('white');
# ax.set_ylim([-2,2])
ax.set_xlabel('x');

```

2 Exercise 2 - Logistic Regression

$$\begin{aligned}
 f(x) &= \sin(x) \\
 x^{(n+1)} &= x^{(n)} - H^{-1} \nabla f(x^{(n)}) \\
 &= x^{(n)} + \frac{\cos(x^{(n)})}{\sin(x^{(n)})}
 \end{aligned}$$

2.1 1.1

```

In [4]: def newton_sin(x0, maxiter = 100):
        i = 0
        stop_cond = True
        while stop_cond:
            x = x0 + np.cos(x0) / np.sin(x0)
            if abs(x - x0) < 1e-4:
                stop_cond = False
            else:
                x0 = x
            if i > maxiter:
                print('fail to converge')
                stop_cond = False
            i += 1
        return x, i
    print(newton_sin(x0 = 1))
    print(newton_sin(x0 = -1))

(1.5707963267948966, 4)
(-1.5707963267948966, 4)

```

2.2 2.2

```
In [5]: data = np.array([[.3, .44, .46, .6], [1, 0, 1, 0]])
w0 = np.array([1., 1.]).T
# print(data.shape)
# define sigmoid function:
sigmoid = lambda x: 1 / (1 + np.exp(-x))

def RLS(y, \
        t, \
        phi, \
        w0, \
        threshold = 1e-10, \
        maxiter = 100, \
        cond = True):

    R = np.diag(y * (1 - y))
    i = 0;
    while cond:
        y = sigmoid(phi.dot(w0))
        tmp = np.multiply(y, (1 - y))
        R = np.diag(tmp);

        Z = phi.dot(w0) - np.linalg.pinv(R).dot(y - t)
        w = np.linalg.inv(phi.T.dot(R.dot(phi))).dot(phi.T.dot(R.dot(Z)))
        if np.sum(abs(w - w0)) < threshold: cond = False;
        if i > maxiter : cond = False
        w0 = w; i += 1
    return w0, i
y, t = data
phi = np.stack((np.ones(y.shape), y))
# print(phi)
# print(phi.shape)
w0, i = RLS(y, t, phi.T, w0)


print('Iterations = {0}, w = {1}'.format(i, w0.T))

Iterations = 7, w = [ 9.78227684 -21.73839298]
```

2.3 2.1

```
In [6]: iris_data = np.loadtxt('iris_data.txt')
x = iris_data[:, :2]; c = iris_data[:, -1]

# set colormaps for labels and probabilities
from matplotlib.colors import ListedColormap
```



test, aklfajklldfjlakdfjklaflijk

```
cmap = sb.color_palette('RdBu', 2)
cmap_class = ListedColormap(cmap)
cmap = sb.color_palette('RdBu', 100)
cmap_prob = ListedColormap(cmap)

fig = figure();
ax = fig.add_subplot(111)
p = ax.scatter(x[:,0 ], x[:,1], c = c, cmap = cmap_class);

ax.set_xlabel('x_1');
ax.set_ylabel('x_2');
cb = fig.colorbar(p); cb.set_label('Class label')
savefig('Figures/21.png')
```

1

Explain why log. regression is a good approach for this data set

2.4 2.2

```
In [7]: w0 = np.array([0,0,0])
phi = np.hstack((np.ones( (x.shape[0],1) ) , x)).T
w, i = RLS(x, c, phi.T, w0)

# sigmoid function
# probability before optimization
prob_classified_before_opt = sigmoid(w0.dot(phi))
```

```

prob_classified_after_opt = sigmoid(w.dot(phi))
c1 = np.mean(prob_classified_after_opt[c == 1])
c2 = 1 - c1
print(('The class probabilities are :'\
      ' class 1 = {0}, class 2 = {1}').\
      format(c1, 1-c1))

print('w0 = {0}, i = {1}'.format(w, i))

```

The class probabilities are : class 1 = 0.49205021284958317, class 2 = 0.5079497871
w0 = [0.00440664 -0.02139153 -0.04930069], i = 3

In [8]: # create feature matrix with bias

```

tmp = np.ones((1, x.shape[0]))
phi = np.vstack([tmp, x.T])

# after optimization
prob_classified_after_opt = sigmoid(w.dot(phi))

# equation 4.90 Bisschop
cross_entropy = lambda y, t:\
    -np.sum( t * np.log(y) + (1 - t) * np.log(1 - y))
# cross entropy with w = [0,0,0]
cross_entropy_before_opt = cross_entropy(prob_classified_before_opt, c)

# cross entropy after optimization using IRS
cross_entropy_after_opt = cross_entropy(prob_classified_after_opt, c)
print('cross-entropy : before = {0}, after = {1} optimization'\
      .format(cross_entropy_before_opt, cross_entropy_after_opt))

# plot before optimization
fig = figure(); fig.clf()
ax = fig.add_subplot(2,1,1)
p = ax.scatter(x[:,0], x[:,1], c = prob_classified_before_opt, cmap = cmap_p)
cb = fig.colorbar(p); cb.set_label('p($c_1 \mid \phi$)')
ax.set_xlabel('x_1')
ax.set_ylabel('x_2')
sb.set_context('poster')
ax.set_title('Before optimization')

# plot after optimization
ax = fig.add_subplot(2,1,2)
p = ax.scatter(x[:,0], x[:,1], c = prob_classified_after_opt, cmap = cmap_p)
cb = fig.colorbar(p); cb.set_label('p($c_1 \mid \phi$)')
ax.set_xlabel('x_1')
ax.set_ylabel('x_2')

```



```

sb.set_context('poster')
ax.set_title('After optimization')

fig.tight_layout()
savefig('Figures/22.png')

```

cross-entropy : before = 693.1471805599454, after = 692.969359482537 optimization

```

In [9]: from scipy.stats import multivariate_normal as mv
        create_basis_vector = lambda mu, sigma: mv(mu, sigma)
        mu1 = np.array([0,0]); mu2 = np.array([1,1])
        sigma2 = .2

        phi1 = create_basis_vector(mu1, sigma2 * np.eye(2))
        phi2 = create_basis_vector(mu2, sigma2 * np.eye(2))

        phi1_pdf = phi1.pdf(x);
        phi2_pdf = phi2.pdf(x);

        phi = np.vstack((np.ones(phi1_pdf.shape), phi1_pdf, phi2_pdf))
        # print(phi.shape, phi1_pdf.shape); assert 0
        w, i = RLS(x, c, phi.T, w0)

        probb_gaussian_after_opt = sigmoid(w.dot(phi))
        cross_entropy_gaussian = cross_entropy(probb_gaussian_after_opt, c)
        print('Cross-entropy with gaussian basis function = {0}'\
              .format(cross_entropy_gaussian))

        fig = figure()
        # subplot: origingal inputspace
        ax = fig.add_subplot(211)
        p = ax.scatter(x[:,0], x[:,1], \
                      c = c, cmap = cmap_class);
        cb = fig.colorbar(p); cb.set_label('Class label')
        ax.set_title('Original input space')
        ax.set_xlabel('$x_1$')
        ax.set_ylabel('$x_2$')

        # basis function decomposition
        ax = fig.add_subplot(212)
        ax.set_xlabel('$x_1$')
        ax.set_ylabel('$x_2$')
        sb.set_context('poster')

        p = ax.scatter(phi1.pdf(x), phi2.pdf(x), \
                      c = c, cmap = cmap_class);

```

```

cb = fig.colorbar(p); cb.set_label('Class label')

ax.set_xlabel('$\phi_1$')
ax.set_ylabel('$\phi_2$')
ax.set_title('basis vector space')
sb.set_context('poster')
fig.tight_layout()
savefig('Figures/24.png')

```

Cross-entropy with gaussian basis function = 346.50408046148465

```
In [10]: fig = figure()
```

```

# optimization results using basis functions
ax = fig.add_subplot(211)
p = ax.scatter(phi1.pdf(x), phi2.pdf(x), \
               c = prob_gaussian_after_opt, \
               cmap = cmap_prob);
cb = fig.colorbar(p); cb.set_label('p($c_1 \mid t)$')
ax.set_xlabel('$\phi_1$')
ax.set_ylabel('$\phi_2$')
ax.set_title('$P(c_1 \mid t)$ in basis space')

ax = fig.add_subplot(212)
p = ax.scatter(x[:,0], x[:,1], \
               c = prob_gaussian_after_opt, \
               cmap = cmap_prob);
cb = fig.colorbar(p); cb.set_label('p($c_1 \mid t)$')

ax.set_xlabel('$x_1$')
ax.set_ylabel('$x_2$')
ax.set_title('$P(c_1 \mid t)$ in original space')
sb.set_context('poster')
fig.tight_layout()
savefig('Figures/24.png')

```